

# **Introduction to Bash**

First Year Bootcamp 2020

# Opening bash

OS X/Linux: Open terminal, bash is default shell

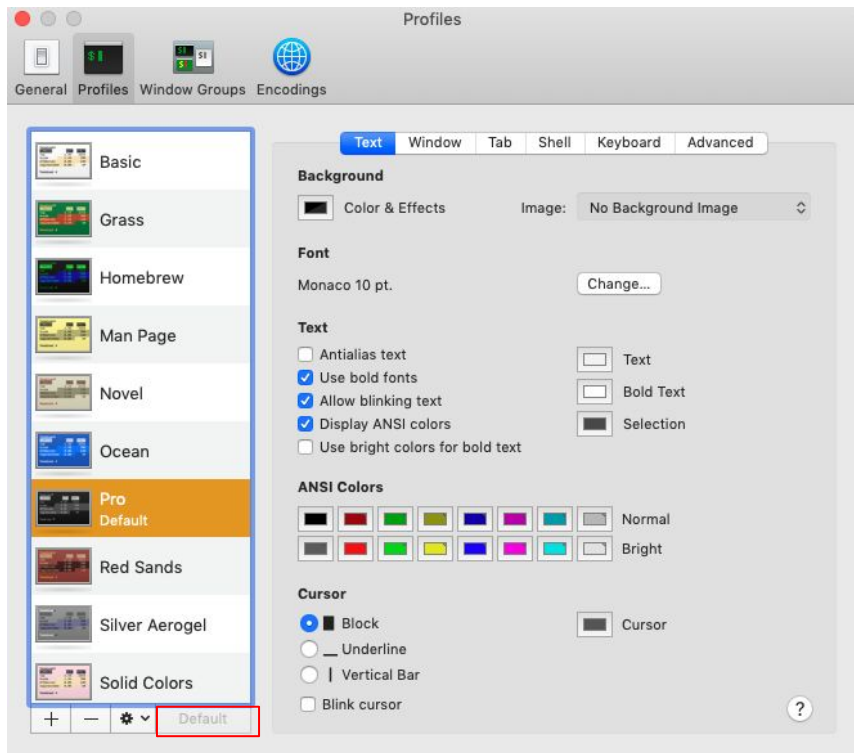
Windows: Go to Start → Git → Git Bash



# Look like a pro

Terminal > Preferences > Text

Select a Profiles tyle you like and click  
“Default”



# There are computers other than mine?

Yes, yes there are. Now we're going to show you how working with files on other computers from bash isn't too much harder than working with files on your own computer, using the tools `ssh` and `scp`. For this example, we'll show you how to use Sherlock.



# Sherlock



Sherlock is a High-Performance Computing (HPC) cluster, operated by the **Stanford Research Computing Center** to provide computing resources to the Stanford community at large.

How much does it cost?

Sherlock is **free** to use for anyone doing sponsored research at Stanford. Any faculty member can request access for research purposes, and get an account with a base storage allocation and unlimited compute time on the global, shared pool of resources. In case those free resources are not sufficient, Stanford Research Computing offers Faculty members the **opportunity to invest** into the cluster, and get access to additional computing resources for their research teams.

Unlike traditional clusters, Sherlock is a collaborative system where the majority of nodes are purchased and shared by the cluster users. When a user (typically a PI) purchases one or more nodes, they become an *owner*.

The resource scheduler configuration works like this:

- owners and their research teams get immediate and exclusive access to the resources they purchased,
- when those nodes are idle, other owners can use them,
- when the purchasing owners want to use their resources, jobs from other owners that may be running on them are preempted (*ie.* killed and re-queued).

# Getting start

To start using Sherlock, you will need:

- an active SUNet ID,
- a Sherlock account,
- a SSH client,
- good understanding of the concepts and terms used throughout that documentation,
- some familiarity with Unix/Linux command-line environments, and notions of shell scripting.

For more information: <https://www.sherlock.stanford.edu/docs/getting-started/prerequisites/>

# Connect to Sherlock

To login to Sherlock, open a terminal and type the following command, where `<sunetid>` should be replaced by your *actual* SUNet ID:

```
$ ssh <sunetid>@login.sherlock.stanford.edu
```

When prompted for your password, type the password that corresponds to your sunet ID. You will need to use two-factor authentication.

Then, you should see a welcome screen:

```
--*-- Stanford Research Computing Center -*--

Sherlock 2.0

-----
This system is for authorized users only and users must comply with all
Stanford computing, network and research policies. All activity may be
recorded for security and monitoring purposes. For more information, see
https://doresearch.stanford.edu/policies/research-policy-handbook and
https://adminguide.stanford.edu/chapter-6/subchapter-2/policy-6-2-1
-----

Sherlock is *NOT* approved for storing or processing HIPAA, PHI, PII nor
any kind of High Risk data. Users are responsible for the compliance of
their data.
See https://uit.stanford.edu/guide/riskclassifications for details.
-----

Docs      https://www.sherlock.stanford.edu/docs
Support   https://www.sherlock.stanford.edu/docs/#support

Web       https://www.sherlock.stanford.edu
News      https://news.sherlock.stanford.edu
Status    https://status.sherlock.stanford.edu
-----
```



# Requesting resources (i.e., submit a job)

Requesting computing resources on Sherlock is done via a resource scheduler, whose very purpose is to match compute resources in the cluster (CPUs, GPUs, memory, ...) with user resource requests.

The scheduler provides three key functions:

1. it allocates access to resources (compute nodes) to users for some duration of time so they can perform work.
2. it provides a framework for starting, executing, and monitoring work (typically a parallel job such as MPI) on a set of allocated nodes.
3. it arbitrates contention for resources by managing a queue of pending jobs

# How to submit a job

A job consists in two parts: resource requests and job steps.

**Resource requests** describe the amount of computing resource (CPUs, GPUs, memory, expected run time, etc.) that the job will need to successfully run.

**Job steps** describe tasks that must be executed.

# Batch scripts

The typical way of creating a job is to write a job submission script. A submission script is a shell script (e.g. a Bash script) whose first comments, if they are prefixed with `#SBATCH`, are interpreted by Slurm as parameters describing resource requests and submissions options.

For instance, the following script would request one task with one CPU for 10 minutes, along with 2 GB of memory, in the default partition:

Features of the  
resources needed  
to run your task

```
#!/bin/bash
#
#SBATCH --job-name=test
#
#SBATCH --time=10:00
#SBATCH --ntasks=1
#SBATCH --cpus-per-task=1
#SBATCH --mem-per-cpu=2G
```

```
srun hostname
srun sleep 60
```

← This is the task you want to  
run on Sherlock

# Job submission and check job

Once the submission script is complete, you can submit the script to the scheduler with the sbatch command:

```
$ sbatch submit.sh  
Submitted batch job 1377
```

Sbatch will return the ID it has assigned to the job

Once submitted, the job enters the queue in the Pending state. You can check the job status using the following command:

↓

```
$ squeue -u $USER
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST(REASON)
1377	normal	test	kilian	R	0:12	1	sh-101-01

↖

Sbatch will return the job status

# Sherlock OnDemand

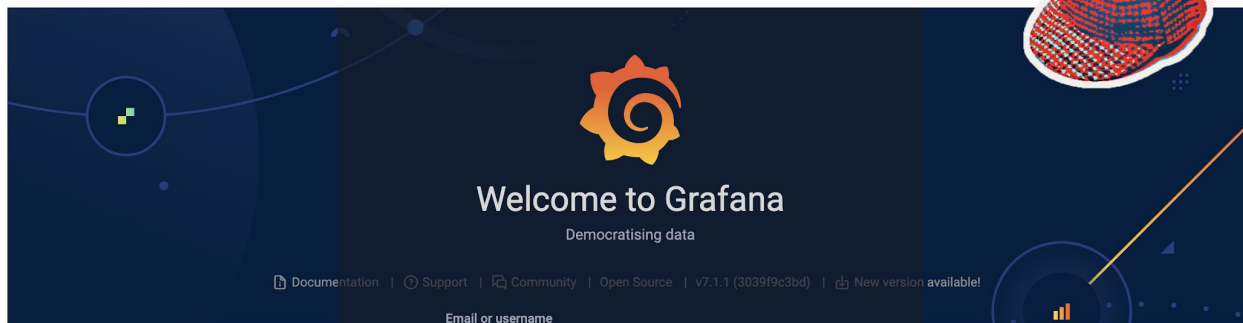
The Sherlock OnDemand interface allows you to conduct your research on Sherlock through a web browser. You can manage files (create, edit and move them), submit and monitor your jobs, see their output, check the status of the job queue, run a Jupyter notebook and much more, without logging in to Sherlock the traditional way, via a SSH terminal connection.

To connect to Sherlock OnDemand, simply point your browser to **<https://login.sherlock.stanford.edu>**

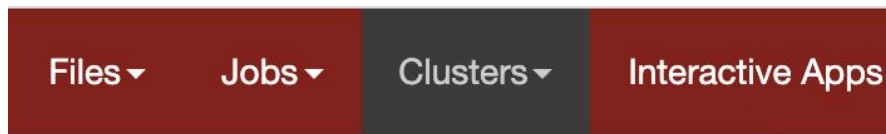
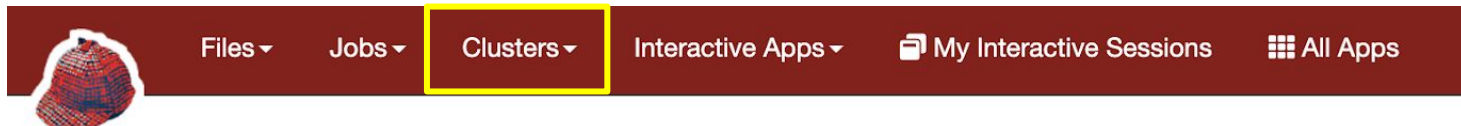
## Welcome to Sherlock OnDemand

An integrated, easy to use, and unified access point for Sherlock resources.

Once you login,  
the home page  
looks like this:

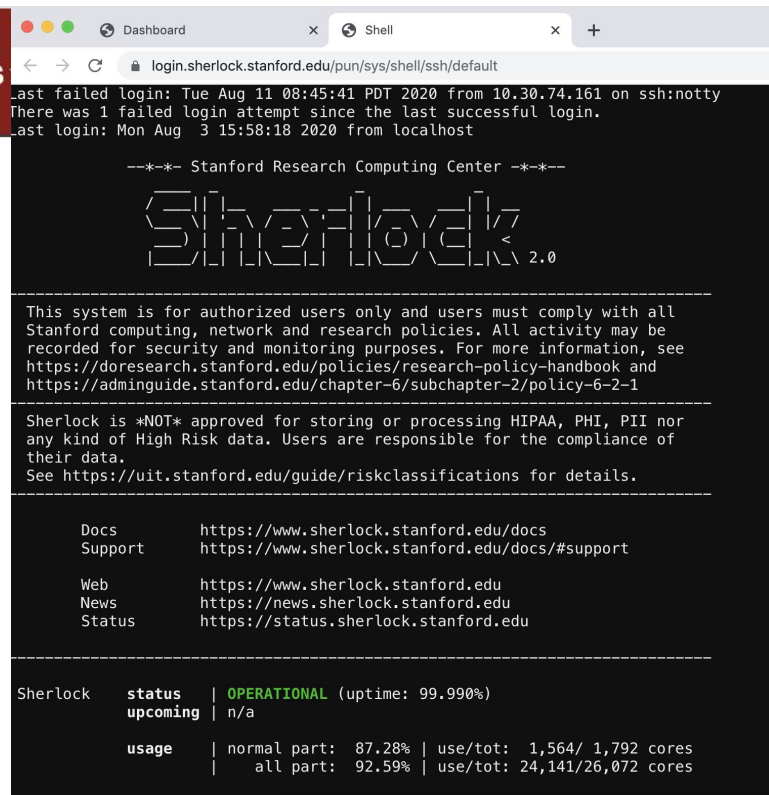


Let's go through a few useful things you can do with Sherlock OnDemand:






>\_Shell Access

You can get shell access to Sherlock in the browser. This is exactly the same as you ssh to Sherlock using terminal.



Interactive Apps ▾My Interactive Sessions

Servers

-  Jupyter Notebook
-  RStudio Server
-  TensorBoard

OnDemand

easy to use, and unified access point for

You can run Jupyter Notebook, RStudio, and TensorBoard on Sherlock using interactive Apps.




For example, you can request to run Jupyter Notebook on Sherlock. Once you click “Jupyter Notebook”, you will need to fill out this form to specify the different parameters for your job (time limit, number of nodes, CPUs, partition to use, etc.).


Once you finish the form, you can click the Launch button to start the session.


Home / My Interactive Sessions / Jupyter Notebook

Interactive Apps

Servers

 Jupyter Notebook

 RStudio Server

 TensorBoard

## Jupyter Notebook

This app will launch a [Jupyter Notebook](#) server on [Sherlock](#).

**Python version**

Python 3.6 ▾

**Extra Jupyter arguments (optional)**

**Additional modules (optional)**

...

Launch

### Interactive Apps

#### Servers

 Jupyter Notebook

 RStudio Server

 TensorBoard

### Jupyter Notebook (32463109)

1 node | 1 core | Running

**Host:** sh-09-06.int

**Created at:** 2018-11-19 16:29:05 PST

**Time Remaining:** about 2 hours

**Session ID:** [9c6d07d4-3a6f-4fa6-975d-e8d337de33cb](#)

 Delete

[Connect to Jupyter](#)

Once your session is ready, you can see this in “My Interactive Sessions”. Click the blue “Connect to Jupyter” button to open Jupyter Notebook!